



CSEN 140L - Lab 9

Nolan Anderson



Lab 9 High-Level Objectives

- Use PCA on the MNIST dataset
- Complete the tasks on the lab assignment PDF a single Jupyter notebook (.ipynb) file
- Use only the **NumPy** and **Matplotlib** libraries for your implementation
 - No machine learning libraries (except for the data retrieval part given by professor)
- Demo to me and submit your notebook to Camino before leaving the lab section



Review: The Data Explained

- The MNIST dataset consists of 28 x 28 grayscale images of handwritten digits (0 – 9).
 - Since the images are 28 x 28, each image has 784 pixels.
 - Each pixel counts as a feature.
 - Because the dataset is quite large, we will just be using the first 2000 samples.
 - As such, your subset of the data should be (2000 x 784).



Principal Component Analysis (PCA)

- Principal Component Analysis is dimensionality reduction method
- The goal is of PCA is preserve the maximum amount of variance by choosing the appropriate axes (i.e. the principal components)
- There are *at most* as many principal components (PCs) as there are features
 - Usually, we want to use enough PCs to at least cover 50% of the variance
 - However, less coverage is okay when plotting since we can only easily interpret plots of 2 or 3 dimensions



The Algorithm (1)

- Prepare the data
 - Load the image data X from the MNIST dataset
 - This should be the first 2000 images
 - Center the data X by subtracting off the mean
- Compute the covariance matrix for all the features
 - Ensure the matrix is $(784, 784)$, transposing if needed
- Find the eigenvectors and eigenvalues of the covariance matrix with `np.linalg.eigh()`
 - You should then transpose your eigenvector result since the columns are the actual eigenvectors



The Algorithm (2)

- Sort the eigenvalues and eigenvectors based on the eigenvalues
 - Should be largest to smallest (descending order)
 - You can find the sorted indices with `np.argsort()`, but you'll need to reverse the order somehow
 - Once sorted, the eigenvectors are the principal components
- Compute how much variance of the data is explained by having 1, 2, and 3 principal components
 - Your result should be between 0 and 1 for each
 - See tips for additional details



The Algorithm (3)

- To actually perform dimensionality reduction, compute the dot product of principal components with the image data
 - For the 2D plot, use the first two PCs for the dot product, result = (784, 2)
 - For the 3D plot, use the first three PCs for the dot product, result = (784, 3)
- Use Matplotlib to construct 2D and 3D scatter plots of the dimensionality-reduced data
 - Use the MNIST target values to have 10 distinct colors
 - Provide a legend to reference which color is which digit
 - See tips for additional details



Tips

- Computing Explained Variance
 - When computing the explained variance, first compute the total sum of all the eigenvalues
 - Then, you can divide the cumulative sum of the eigenvalues by the total to get an array of explained variance
 - You can also use a sum variable and a for loop to do a similar computation, your choice
 - Since we divided a partial sum by the total, we get values between 0 and 1.



Tips

- Plotting in Color
 - To get the target labels (the digits), use `mnist["target"]`.
 - Then, I suggest you construct a numpy array with those labels with `dtype=int` since `scatter` expects an array of integers, not strings
 - To actually implement the different colors, use the `c` field of `scatter()` and set that equal to your array for the digits. Then use the `cmap` field to specify a color map string (e.g. "gist_rainbow")
 - For more details, consult the following links:
 - [pyplot.scatter](#) | [colormap reference](#)



Lab Outcomes | What I Need To See


- Print out the amount of variance covered by having 1, 2, and 3 principal components (three values)
- Display two plots: a 2D plot and 3D plot
 - Your plots must have legends and different colors for different digits



Helpful NumPy Functions

- `np.linalg.eigh()`
 - Compute the eigenvalues and eigenvectors
- `np.argsort()`
 - Get sorted indices in ascending order
- `np.cumsum()`
 - Cumulative sums of an array
- `np.mean()`
- `np.sum()`
- `np.dot()`

Sample 2D Plot



Sample 3D Plot

